

Problem Set 2

1. The exponential distribution

Problem a

The support of X is $x \geq 0$

Problem b

$$L(\theta) = \prod_{i=1}^n \theta e^{-\theta x_i} \log L(\theta|x) = n \log \theta - \theta \sum_{i=1}^n x_i$$

Problem c

Let $\frac{d}{d\theta} \log L(\theta) = \frac{n}{\theta} - \sum_{i=1}^n x_i = 0$ We have $\hat{\theta} = \frac{n}{\sum_{i=1}^n x_i}$

Problem d

```
library(formatR)
```

Warning: package 'formatR' was built under R version 4.4.2

```
set.seed(5)
x <- rexp(10000, 5)


exp.ll = function(theta, x){
  n = length(x)
  return(n*log(theta) - theta*sum(x))
}
```

```

theta_seq = seq(1, 10, length.out = 100)
log_log_vals = sapply(theta_seq, function(theta) exp.ll(theta, x))

plot(theta_seq, log_log_vals, type="l",
      xlab="theta", ylab="Likelihood", col="purple", lwd=2)

```



It looks like around 5

```

theta_tilde = 6
theta_hat = length(x)/sum(x)

ll_hat = exp.ll(theta_hat, x)
ll_tilde = exp.ll(theta_tilde, x)

likelihood_ratio = exp(exp.ll(theta_tilde, x) - exp.ll(theta_hat, x))

print(likelihood_ratio)

```

```
[1] 4.866226e-81
```

```
log_likelihood = function(theta, x) {
  n = length(x)
  return(- (n*log(theta)-theta*sum(x)))
}
```

```
start_time_BFGS = Sys.time()
result_BFGS = optim(par=1, fn=log_likelihood, x=x, method="BFGS")
```

```
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
```

```
end_time_BFGS = Sys.time()
time_BFGS = end_time_BFGS - start_time_BFGS
start_time_SANN = Sys.time()
result_SANN = optim(par=1, fn=log_likelihood, x=x, method="SANN")
end_time_SANN = Sys.time()
time_SANN = end_time_SANN - start_time_SANN

print(result_BFGS$par)
```

```
[1] 4.980155
```

```
print(result_SANN$par)
```

```
[1] 4.979926
```

```
print(time_BFGS)
```

```
Time difference of 0.004085064 secs
```


```
print(time_SANN)
```

```
Time difference of 0.07587695 secs
```

2. Maximizing a multivariate function

Problem a

```
mvn <- function(xy) {  
  x <- xy[1]  
  y <- xy[2]  
  z <- exp(-0.5 * ((x - 2)^2 + (y - 1)^2))  
  return(z)  
}  
  
# install.packages('lattice')  
library(lattice)  
  
y <- x <- seq(-5, 5, by = 0.1)  
grid <- expand.grid(x, y)  
names(grid) <- c("x", "y")  
grid$z <- apply(grid, 1, mvn)  
  
wireframe(z ~ x + y, data = grid, shade = TRUE, light.source = c(10, 0, 10), scales = list(a
```



It looks like when $x = 0, y = 0$ the function achieves a maximum.


Problem b

```
neg_mvn <- function(xy) {  
  return(-mvn(xy))  
}  
  
opt1 <- optim(c(1, 0), neg_mvn, method = "BFGS")  
opt2 <- optim(c(5, 5), neg_mvn, method = "BFGS")  
  
data.frame(  
  Start_Point = c("(1,0)", "(5,5)"),  
  Optimum_X = c(opt1$par[1], opt2$par[1]),  
  Optimum_Y = c(opt1$par[2], opt2$par[2]),  
  Function_Value = c(-opt1$value, -opt2$value)  
)
```

	Start_Point	Optimum_X	Optimum_Y	Function_Value
1	(1,0)	2.000000	1.000000	1.000000e+00
2	(5,5)	4.998888	4.998518	3.761331e-06

The optimization starting from (5,5) did not converge to (2,1). The function value is very small, meaning that it has not reached the peak.

```
y <- x <- seq(-5, 5, by = 0.1)  
grid <- expand.grid(x, y)  
names(grid) <- c("x", "y")  
grid$log_z <- apply(grid, 1, mvn)  
wireframe(log_z ~ x + y, data = grid, shade = TRUE, light.source = c(10, 0, 10), scales = li
```



Problem c

```
neg_log_mvn <- function(xy) {
  return(-mvn(xy))
}

opt1_log <- optim(c(1, 0), neg_log_mvn, method = "BFGS")

opt2_log <- optim(c(5, 5), neg_log_mvn, method = "BFGS")

data.frame(
  Start_Point = c("(1,0)", "(5,5)"),
  Optimum_X = c(opt1_log$par[1], opt2_log$par[1]),
  Optimum_Y = c(opt1_log$par[2], opt2_log$par[2]),
  Log_Function_Value = c(-opt1_log$value, -opt2_log$value)
)
```

	Start_Point	Optimum_X	Optimum_Y	Log_Function_Value
1	(1,0)	2.000000	1.000000	1.000000e+00
2	(5,5)	4.998888	4.998518	3.761331e-06

Using log-likelihood improves numerical stability by preventing underflow and helps the optimizer converge faster by smoothing sharp variations. While function values change due to the logarithm, the optimal (x, y) remains the same.

3. The normal variance

Problem a

log normal distribution is:

$$\ell(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2.$$

derivative:

$$\frac{\partial \ell}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \mu)^2.$$

equal to zero:

$$-\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \mu)^2 = 0.$$

then

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2$$

Problem b

- The **MLE for variance** is $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$
- The **MLE is biased**, as it systematically underestimates σ^2
- The **unbiased estimator** is $s^2 = \frac{n}{n-1} \hat{\sigma}^2$

Appendix

I certify that we did not use any LLM or generative AI tool in this assignment.