

# Problem Set 2

## 1. The exponential distribution

### Problem a

The support of  $X$  is  $x \geq 0$

### Problem b

$$L(\theta) = \prod_{i=1}^n \theta e^{-\theta x_i} \log L(\theta|x) = n \log \theta - \theta \sum_{i=1}^n x_i$$

### Problem c

Let  $\frac{d}{d\theta} \log L(\theta) = \frac{n}{\theta} - \sum_{i=1}^n x_i = 0$  We have  $\hat{\theta} = \frac{n}{\sum_{i=1}^n x_i}$

### Problem d

```
library(formatR)
```

Warning: package 'formatR' was built under R version 4.4.2

```
set.seed(5)
x <- rexp(10000, 5)


exp.ll = function(theta, x){
  n = length(x)
  return(n*log(theta) - theta*sum(x))
}
```

```

theta_seq = seq(1, 10, length.out = 100)
log_log_vals = sapply(theta_seq, exp.ll, x = x)

plot(theta_seq, log_log_vals, type="l",
      xlab="theta", ylab="Likelihood", col="purple", lwd=2)

```



It looks like around 5

```

theta_tilde = 6
theta_hat = length(x) / sum(x)

ll_hat = exp.ll(theta_hat, x)
ll_tilde = exp.ll(theta_tilde, x)

likelihood_ratio = exp(ll_hat - ll_tilde)
print(likelihood_ratio)

```

```
[1] 2.05498e+80
```

```

log_likelihood = function(theta, x) {
  n = length(x)

```

```
    return(- (n*log(theta)-theta*sum(x)))
}
```

```
start_time_BFGS = Sys.time()
result_BFGS = optim(par=1, fn=log_likelihood, x=x, method="BFGS")
```

```
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
Warning in log(theta): NaNs produced
```

```
end_time_BFGS = Sys.time()
time_BFGS = end_time_BFGS - start_time_BFGS
start_time_SANN = Sys.time()
result_SANN = optim(par=1, fn=log_likelihood, x=x, method="SANN")
end_time_SANN = Sys.time()
time_SANN = end_time_SANN - start_time_SANN

print(result_BFGS$par)
```

```
[1] 4.980155
```

```
print(result_SANN$par)
```

```
[1] 4.979926
```

```
print(time_BFGS)
```

```
Time difference of 0.003666878 secs
```


```
print(time_SANN)
```

```
Time difference of 0.07271194 secs
```

## 2. Maximizing a multivariate function

### Problem a

```
mvn <- function(xy) {  
  x <- xy[1]  
  y <- xy[2]  
  z <- exp(-0.5 * ((x - 2)^2 + (y - 1)^2))  
  return(z)  
}  
  
# install.packages('lattice')  
library(lattice)  
  
y <- x <- seq(-5, 5, by = 0.1)  
grid <- expand.grid(x, y)  
names(grid) <- c("x", "y")  
grid$z <- apply(grid, 1, mvn)  
  
wireframe(z ~ x + y, data = grid, shade = TRUE, light.source = c(10, 0, 10), scales = list(a
```



It looks like when  $x = 0, y = 0$  the function achieves a maximum.


## Problem b

```
neg_mvn <- function(xy) {  
  return(-mvn(xy))  
}  
  
opt1 <- optim(c(1, 0), neg_mvn, method = "BFGS")  
opt2 <- optim(c(5, 5), neg_mvn, method = "BFGS")  
  
data.frame(  
  Start_Point = c("(1,0)", "(5,5)"),  
  Optimum_X = c(opt1$par[1], opt2$par[1]),  
  Optimum_Y = c(opt1$par[2], opt2$par[2]),  
  Function_Value = c(-opt1$value, -opt2$value)  
)
```

	Start_Point	Optimum_X	Optimum_Y	Function_Value
1	(1,0)	2.000000	1.000000	1.000000e+00
2	(5,5)	4.998888	4.998518	3.761331e-06

The optimization starting from (5,5) did not converge to (2,1). The function value is very small, meaning that it has not reached the peak.

```
y <- x <- seq(-5, 5, by = 0.1)  
grid <- expand.grid(x, y)  
names(grid) <- c("x", "y")  
grid$log_z <- apply(grid, 1, mvn)  
wireframe(log_z ~ x + y, data = grid, shade = TRUE, light.source = c(10, 0, 10), scales = li
```



### Problem c

```
neg_log_mvn <- function(xy) {
  return(-mvn(xy))
}

opt1_log <- optim(c(1, 0), neg_log_mvn, method = "BFGS")

opt2_log <- optim(c(5, 5), neg_log_mvn, method = "BFGS")

data.frame(
  Start_Point = c("(1,0)", "(5,5)"),
  Optimum_X = c(opt1_log$par[1], opt2_log$par[1]),
  Optimum_Y = c(opt1_log$par[2], opt2_log$par[2]),
  Log_Function_Value = c(-opt1_log$value, -opt2_log$value)
)
```

	Start_Point	Optimum_X	Optimum_Y	Log_Function_Value
1	(1,0)	2.000000	1.000000	1.000000e+00
2	(5,5)	4.998888	4.998518	3.761331e-06

Using log-likelihood improves numerical stability by preventing underflow and helps the optimizer converge faster by smoothing sharp variations. While function values change due to the logarithm, the optimal  $(x, y)$  remains the same.

### 3. The normal variance

#### Problem a

log normal distribution is:

$$\ell(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2.$$

derivative:

$$\frac{\partial \ell}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \mu)^2.$$

equal to zero:

$$-\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \mu)^2 = 0.$$

then

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2$$

#### Problem b

##### Step 1: MLE for Variance

From Problem a, the maximum likelihood estimator (MLE) for variance is  $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$

##### Step 2: Expected Value of $\hat{\sigma}^2$

To check if  $\hat{\sigma}^2$  is biased, take its expectation:  $E[\hat{\sigma}^2] = E\left[\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2\right]$

By linearity of expectation:  $E\left[\sum_{i=1}^n (x_i - \mu)^2\right] = \sum_{i=1}^n E[(x_i - \mu)^2]$

Since  $x_i \sim \mathcal{N}(\mu, \sigma^2)$ , we know  $E[(x_i - \mu)^2] = \sigma^2$

Thus,  $E[\hat{\sigma}^2] = \frac{1}{n} \cdot n\sigma^2 = \sigma^2$

This suggests that  $\hat{\sigma}^2$  is an unbiased estimator **only if**  $\mu$  is known.

### Step 3: Why is the MLE Biased?

In practice, we do not know  $\mu$  and instead estimate it with the **sample mean**:  $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$

Using  $\bar{x}$  instead of  $\mu$ , the variance estimator becomes  $S^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$

Taking expectation:  $E[S^2] = \frac{1}{n} \sum_{i=1}^n E[(x_i - \bar{x})^2]$

It can be shown that:  $E[(x_i - \bar{x})^2] = \frac{n-1}{n} \sigma^2$

which leads to:  $E[S^2] = \frac{n-1}{n} \sigma^2$

Since  $E[S^2] \neq \sigma^2$ , the MLE underestimates  $\sigma^2$  by a factor of  $\frac{n-1}{n}$ .

### Step 4: Unbiased Estimator

To correct the bias, we rescale  $S^2$ :  $s^2 = \frac{n}{n-1} \hat{\sigma}^2$

This gives the **unbiased variance estimator**:  $s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$

which satisfies:  $E[s^2] = \sigma^2$

### Final Answer:

- **MLE for variance:**

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

- **MLE is biased**, underestimating  $\sigma^2$  by a factor of  $\frac{n-1}{n}$ .
- **Unbiased estimator**:  $s^2 = \frac{n}{n-1} \hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$

This adjustment ensures an unbiased estimate of  $\sigma^2$ .

### Appendix

I certify that we did not use any LLM or generative AI tool in this assignment.